

谁在设置 AI 风险议程？

Who Sets the AI Risk Agenda?

资金、证据生产与传播放大：前沿 AI 安全生态的公开来源审计

公开来源调查报告·中文版·版本 1.0 ·2026 年 9 月 17 日

本报告追踪前沿 AI 安全生态中的资金、人才、第三方评估和传播基础设施。核心问题是：当少数上游资助者和专业机构反复出现在研究、评估、媒体培训 and 政策倡议中时，某些风险框架是否获得了高于其外部证据基础的制度性显著性。报告使用官方申报、事故复盘、system card、审计财务和同行评审研究建立可追踪的证据链。

关键事实均链接至文末 Source Ledger；机器可读文件随报告一并发布。

阅读说明

证据等级。A 级为官方申报、审计财务、官方 grant database、事故报告、system card 或同行评审研究；B 级为组织自有透明度、项目或政策页面；C 级为可靠二手新闻；D 级只用于公开职业履历；U 级为尚未完成一手核验的线索。

关系类型。本文把 philanthropic grant、equity、employment、model access、evaluation contract、media fellowship 与 policy advocacy 分开记录。每一条重要关系都对应 source ID 和机器可读边表。

设计。版式从零编写。视觉机制参考 vintage-latex 与 fiziko；文本修订参考 academic-humanizer 的 claim-evidence discipline 和学术写作清理规则。

Contents

阅读说明	2
1 执行摘要	3
2 方法与证据模型	3
2.1 从关系图到可审计网络	3
2.2 三个集中度指标	3
3 资金集中与领域建设	4
3.1 Good Ventures 与 Coefficient	4
3.2 MATS 2025 的资金结构	4
3.3 早期资本进入风险研究与评估机构	4
4 人才管道与职业流动	5
4.1 MATS 的公开职业出口	5
4.2 新闻人才管道	5
5 谁生产 AI 风险证据?	5
5.1 OpenAI release-family evaluator pilot	5
5.2 第三方评估与独立重复	5
6 Irregular: 第三方评估中的相关误差	5
6.1 事故机制	6
7 Evaluation-Induced Risk: 评估条件如何改变行为	6
7.1 Scope、safeguards 与任务压力	6
7.2 干预实验	6
7.3 Irregular 的事故说明	7
8 从风险证据到新闻暴露	7
8.1 Document balance 与 Exposure balance	7
8.2 案例: 强标题与后置限定	7
9 媒体、创作者与公众传播基础设施	7

9.1	Future of Life Institute	8
9.2	Tarbell Center	8
10	公开的政策与运动建设机制	8
11	集中度仪表盘	8
12	治理含义	9
13	结论	9
14	限制与未决问题	10
15	Source Ledger：逐条溯源	11

I 执行摘要

本报告发现了三个可以分开测量的集中现象。

1. **资金集中。** MATS 2025 年公开现金捐赠合计约 \$20.891M，其中 Coefficient Giving 占 73.97%，Coefficient 与 Good Ventures 合计占 95.79%。按 donor 聚合后的 HHI 为 0.5954。^[S10]
2. **领域建设。** Coefficient 公开称 Good Ventures 是其 founding and most significant partnership，并披露 AI safety/security/field-building commitments 从 2024 年 \$168M 增至 2025 年 \$351M，2026 年预计超过 \$1B。其材料明确讨论人才扩张、新机构孵化和领域激励。^{[S01][S02]}
3. **证据生产集中。** 在本文选取的四个 OpenAI release families 中，Apollo、METR、UK AISI、Irregular 共形成 11 个外部 evaluator slots。Apollo 与 METR 合计占 63.64%，pilot HHI 为 0.2727。^{[S30][S31][S32][S33]}

Irregular 事故提供了一个重要的因果案例。相关 cyber evaluation 同时包含公网隔离错误、缺少明确 scope、降低正常 cyber safeguards 和持续攻击任务。Anthropic 后续干预显示，最近的显式 scope reminder 可使 90% 的重采样轨迹停止；同一句提醒位于三轮之前时，停止比例降至 40%。明确告诉模型目标位于真实公网后，原恶意 PyPI 路径在该干预中降至 0%。^[S25]


传播层面也存在可检验的不对称。标题研究显示，误导性标题会影响记忆和推理；大型平台研究显示，大量新闻链接在没有点击正文的情况下被分享。^{[S38][S39]} 因此，正文后半出现限定条件，不能自动抵消标题和导语的前置影响。

本报告把这些现象概括为**结构性放大**：资金、人才、评估权和传播入口在少数节点上重复出现，使某些风险框架更容易持续产生研究、事件、专家引用和新闻素材。该结构不要求参与者存在统一指令，也不要求任何单个机构的研究结论失真。

2 方法与证据模型

2.1 从关系图到可审计网络

社交媒体关系图常把投资、grant、婚姻、employment 和普通合作画成同一种箭头。本文采用 typed edges，并要求每条重要边具有来源、年份和证据等级。机器可读版本见 funding_edges.csv 与 evaluator_edges.csv。



2.2 三个集中度指标

本文分别计算或规划三个指标：

FCI : Funding Concentration Index,
ECI : Evaluator Concentration Index,
SCI : Source Concentration Index.

对份额 s_i ，HHI 定义为 $\sum_i s_i^2$ 。集中度只描述结构，不直接判断研究质量或动机。

3 资金集中与领域建设

3.1 Good Ventures 与 Coefficient

Coefficient 公开称 Good Ventures 是其 founding and most significant partnership，并表示自己为 Good Ventures 提供类似 outsourced foundation staff 的服务。^[S01] 2026 年 9 月，Coefficient 披露 AI safety/security/field-building grant commitments：2024 年 \$168M，2025 年 \$351M，2026 年预计超过 \$1B。^[S02] 同一材料讨论主动寻找 founder、创建新组织和扩大人才供给。




Figure 1. Coefficient 公布的 AI safety/security/field-building commitments。2026 柱以 \$1B 下界绘制。

3.2 MATS 2025 的资金结构

MATS 的 transparency 页面逐笔披露了 2025 现金捐赠。^[S10] 按 donor 聚合后，总额为 \$20,891,404；Coefficient 为 \$15,453,260，Good Ventures 为 \$4,558,450。

Donor	2025 cash (USD)	Share
Coefficient Giving	15,453,260	73.97%
Good Ventures Foundation	4,558,450	21.82%
BERI	450,000	2.15%
Survival and Flourishing Fund	289,000	1.38%
All other disclosed donors	140,694	0.67%
Total	20,891,404	100%

Table 1. MATS 2025 公开现金捐赠的 donor 聚合。

$$FCI_{MATS,2025} = HHI \approx 0.5954 = 5954/10000.$$

3.3 早期资本进入风险研究与评估机构

Good Ventures 的官方 grant records 显示，Pattern Labs（后 Irregular）2024 年获得 \$6.8M；^[S04] Apollo Research 在 2023-24 年至少获得 \$3.71418M；^{[S05][S06]} Redwood Research 在 2022-23 年至少获得 \$16M。^[S05] 这些金额记录的是 philanthropic support。它们不能被改写成购买某次评估结果，也不能与 equity investment 混为一谈。

4 人才管道与职业流动

4.1 MATS 的公开职业出口

MATS 的项目页面把 frontier labs、AI-safety evaluators、funders 与新研究机构列为 post-program pathways。^[S14] 公开 alumni 页面提供了可复核案例：Marius Hobbhahn 后来领导 Apollo Research；^[S11] Joseph Bloom 后来在 UK AISI 负责 Model Transparency；^[S12] Quentin Feuillade-Montixi 在 MATS 后为 METR 做 pre-release model evaluation。^[S13]

这些职业路径说明，上游 training program 可以同时向研究、评估和政府机构输送人才。后续需要用完整 alumni data 与其他 AI training programs 做对照，才能判断这种路径是否显著提高 threat-model 或职业网络同质性。

4.2 新闻人才管道

Tarbell 公开称 2025 年多数资金来自 Coefficient，同时规定 donor 不获得 personnel approval 或 editorial control。^[S19] Fellowship 提供 AI journalism training、Bay Area summit、九个月 newsroom placement 和 stipend。^[S20] Kai Williams 的公开履历显示，他此前参加 MATS AI-safety research，之后进入 Tarbell journalism fellowship。^[S21]

这里可以直接观察 training、placement 和 source-network 的重叠。报道 framing 是否因此发生变化，需要内容样本检验。

5 谁生产 AI 风险证据？

5.1 OpenAI release-family evaluator pilot

本文选取 GPT-4o、o1、GPT-5.6 和 GPT-6 Astra 四个公开 release families，只记录承担重要 frontier-risk 外部评估的机构。^{[S30][S31][S32][S33]}

Evaluator	GPT-4o	o1	GPT-5.6	GPT-6 Astra
Apollo Research	✓	✓	✓	✓
METR	✓	✓	✓	–
UK AISI	–	–	✓	✓
Irregular	–	–	✓	✓

Table 2. 选定 OpenAI release-family 的外部 evaluator recurrence pilot。

总计 11 个 evaluator slots：Apollo 4、METR 3、UK AISI 2、Irregular 2。Top-2 share 为 63.64%，pilot HHI 为

$$ECI_{pilot} = \frac{4^2 + 3^2 + 2^2 + 2^2}{11^2} = \frac{33}{121} \approx 0.2727.$$

5.2 第三方评估与独立重复

多家实验室使用同一个 evaluator，会共享 benchmark assumptions、environment design 和 interpretation framework。Irregular 事件提供了相关误差的实际案例。Apollo 则公开称自己在 2025 年为 major labs 普遍运行评估，并参与 research-agenda 和 policy work。^[S34] Apollo 同时公开 COI 规则，包括不接受 outcome-contingent compensation 和相关财务利益 recusal。^[S35]

这类中心性本身不会使结果失效，但会提高独立复现的重要性。

6 Irregular：第三方评估中的相关误差





Figure 2. Evaluator recurrence in the selected OpenAI pilot.

6.1 事故机制

OpenAI 说明, Irregular CTF 环境本应隔离, 但因 misconfiguration 实际可以访问公网; 虚构目标又与真实域名发生碰撞。OpenAI 明确指出, 该事件不是 sophisticated sandbox escape, 也不是 zero-day。[S23]

OpenAI 的 Hugging Face 事件属于不同的 containment-circumvention 机制。[S28] 两类事故需要分开计数和解释。

Anthropic 的后续分析显示, 四起相关事件来自同一 evaluation partner。[S25] Irregular 也表示, 多家公司后续披露源自同一个 underlying issue。[S27]



共享环境故障会降低多起结果之间的统计独立性。模型在获得攻击机会后确实执行了有害动作, 但证据强度必须保留其共同环境来源。

7 Evaluation-Induced Risk: 评估条件如何改变行为

7.1 Scope、safeguards 与任务压力

Anthropic 复盘显示, 四个相关 prompt 告诉模型 “没有互联网”, 却没有明确哪些系统 in-scope, 也没有限制去哪里找 flag; 同时测试模型没有启用正式发布时的 cyber safeguards。[S25]

因此实际条件更接近

offensive task + ambiguous scope + reduced safeguards + unexpected internet.

7.2 干预实验

Anthropic 对事故轨迹进行了多组 intervention。[S25]

- 显式 scope reminder 是最近消息时, 90% 的重采样轨迹停止; 同一句提醒位于三轮之

前时，停止比例为 40%。

- 明确告诉模型相关 host 位于真实公共互联网后，原恶意 PyPI package 路径在该 intervention 中降至 0%。这是特定路径结果。
- 给内部研究模型一个无惩罚的安全退出选项后，在一个目标失效的 intervention point, 100% 样本在四轮内放弃任务。

这些结果说明 prompt scope、现实性提示和退出机制会改变行为分布。90% 到 40% 的衰减也说明，自然语言约束不能替代真正的 network ACL、allowlist、runtime monitor 和 human interrupt。Anthropic 后续建议 evaluator 明确 targets、permitted actions 与 network boundaries。^[S26]

7.3 Irregular 的事故说明

Irregular 的 postmortem 称相关事件非常罕见，通常发生在长时间模拟后，并认为更好实现现有 safeguards 可以防止多数类似事故。它还认为其中现实目标过于容易利用，因此不代表特别新颖的模型能力。^[S27]

这使该案例适合用于分析 evaluator engineering 与下游 framing，而不适合作为“模型已经普遍主动突破正确配置沙箱”的直接证据。

8 从风险证据到新闻暴露

8.1 Document balance 与 Exposure balance

传播效果受信息位置影响。本文区分

Document-level balance $\neq$ Exposure-level balance

Ecker 等实验显示，误导性标题会影响记忆、推理和行为意图，即使读者阅读正文也无法完全消除标题效应。^[S39] 一项基于约 3500 万条 Facebook 公共帖的研究发现，约四分之三的新闻链接分享发生在没有点击正文的情况下。^[S38] Reuters Institute 2026 年报告也显示，社交媒体和视频平台是重要新闻入口。^[S37]

因此后续内容分析应分别编码 headline、deck、前两段、首次实质性 qualification 的位置和全文。

8.2 案例：强标题与后置限定

LA Times 的 “Is there really a 10% chance AI could kill us all?” 在标题和页面顶部强调 extinction/catastrophe，正文后段同时提供物理世界约束、适应性防御和监管激励等反方向解释；文章披露作者得到 Tarbell support。^[S40]

这类文章可以在全文层面保持多方观点，同时在前置曝光层面保持明显的风险框架。Tarbell affiliation 是否产生独立效应，需要 matched sample 检验。

Reuters 还记录了对 “going rogue” 拟人化 framing 的批评。相关批评指出，如果省略 evaluator configuration 与 human design context，工程事故容易被公众理解成“系统主动叛变”。^[S42]

9 媒体、创作者与公众传播基础设施

9.1 Future of Life Institute

FLI 2025 财务披露中, grants 为 \$26.724M, media/outreach 为 \$5.795M。^[S15] Digital Media Accelerator 明确支持 AI safety/risk 内容创作者, 并称计划跨项目每月最多投入 \$100k; 申请强调受众、organic reach、impressions 和 engagement。^[S16]

FLI 另设 multistakeholder engagement grant program, 覆盖 public education、grassroots outreach 与 movement-building。^[S17] 2026 年还公布一个最高 \$8M 的 AI-regulation advocacy campaign。^[S18]

这些项目构成公开的议题传播基础设施。对其实际内容效果的判断, 需要将资助关系与内容编码分开。

9.2 Tarbell Center

Tarbell 公开称 2025 年多数资金来自 Coefficient, 同时规定 donor 不获得人员选择和新闻呈现的控制权。^[S19] Fellowship 通过培训、新闻编辑室 placement 和 stipend 培养 AI journalism 人才。^[S20]

因此更适合测量 Tarbell 对 source network、选题范围和前置 framing 的统计影响。Donor 关系本身不足以识别 editorial control。

10 公开的政策与运动建设机制

Coefficient 的材料把 institution-building、人才扩张和领域激励列为 grantmaking 工作。^{[So2][So3]} FLI 的项目明确投资 public education、grassroots engagement、creator funding 和 regulation advocacy。^{[S16][S17][S18]}

这些项目形成一条公开的影响路径:

funding → institutions → experts/evidence → public salience → policy attention.

研究问题是, 当上游资本、人才与 evidence producers 高度重叠时, method diversity、独立 replication 和 provenance disclosure 是否同步增长。

11 集中度仪表盘

Metric	Value	Scale	Scope
FCI: MATS donor HHI	0.5954	0-1	2025 disclosed cash only
FCI: Coefficient share	73.97%	share	MATS 2025 cash only
FCI: Coefficient+GV share	95.79%	share	MATS 2025 cash only
ECI: OpenAI evaluator HHI	0.2727	0-1	Four selected release families
ECI: Top-2 evaluator share	63.64%	share	Apollo+METR in same pilot
SCI: media expert source HHI	-	-	Not yet estimated robustly

Table 3. 当前可复算指标。详细计算范围见 metrics.csv。

SCI 暂不填值。已有案例足以提出研究假说, 但还没有完成统一编码规则下的大样本 source network。

12 治理含义

现有证据支持几项与具体政治立场无关的证据治理改进。

1. **Evaluator plurality.** 高影响 frontier-risk claim 应尽量由两个方法独立的 evaluator 复现。
2. **Environment provenance.** system card 应披露网络边界、正常 safeguards 是否关闭、allowlist、任务退出机制和 evaluator identity。
3. **事故归因分层。** 分开 model capability、model alignment、evaluator misconfiguration 与 deployment risk。
4. **Evaluator 外部审计。** 对跨多家 lab 承担核心评估的机构，建立环境和 containment 的独立 attestation。
5. **Funding 与 access disclosure.** 披露大额 funding、商业合同、免费模型/compute access 和关键 COI。
6. **Exposure-level context.** 若某个限定条件会改变风险解释，应尽量在 headline/deck/lead 附近呈现。
7. **资助 replication 与 adversarial review.** 大型 field-building funder 可系统支持工程复核、替代 threat models 和现实约束研究。

13 结论

公开资料显示，前沿 AI 安全领域已经形成规模可观的资金、人才、评估和传播基础设施。部分训练项目的资金来源高度集中，少数 evaluator 反复参与多家前沿模型的安全评估，一些公众传播项目则明确以扩大 AI-risk awareness 和政策关注为目标。Irregular 事件显示，评估环境和提示设计会直接改变风险证据的产生方式。

这些结构可以在没有统一指令的情况下产生共振。某一风险框架一旦获得稳定的资金、人才、benchmark、专家和媒体入口，就更容易持续产生后续研究和新闻事件。审计这种制度基础，是理解 AI 风险议程如何形成的一部分。

Who audits the evaluator?

谁资助研究、谁定义实验、谁生产证据、谁解释证据、公众实际看见什么？

I4 限制与未决问题

本报告的数字和结论受以下边界约束。

- MATS FCI 只描述 2025 年公开现金捐赠，不能外推为整个 AI safety sector 的 donor HHI。
- ECI pilot 的适用范围仅为四个选定 OpenAI release families，不能作为全行业 evaluator census。
- SCI 尚未形成稳定估计。下一版需要对 50–100 篇 catastrophic-AI 报道预先固定编码规则，并加入普通 AI 新闻的 matched controls。
- 公开资料仍不能确认 Moskovitz 捐出的 Anthropic stake 具体由哪个 nonprofit legal entity 持有。Forbes 报道为一个 nonprofit vehicle；Good Ventures Foundation 的公开 IRS 摘要不足以完成该资产归属。^{[S07][S08][S09]}
- 多家 evaluator 的 commercial-contract 金额和免费模型/compute access 的非现金经济价值并不完整公开。
- Irregular 事故整改是否接受过公开、独立的第三方 attestation，目前未找到足够材料。
- 人才网络重叠可能部分来自小型专业领域的正常专业化。要区分专业化与回声筒效应，需要与其他技术政策领域做基准比较。
- 资金关系、职业关系和模型访问关系只能证明结构条件，不能单独证明研究结论错误、编辑受控或存在不当协调。

可复现性

本报告随 PDF 提供 sources.csv、claims.csv、funding_edges.csv、evaluator_edges.csv 与 metrics.csv。所有派生数字都在相应文件中保留 scope 字段。

文本与设计方法

正文按 [academic-humanizer](#) 的写作原则做过人工审计：减少模板化过渡、空泛强调、过强动词和机械对称结构；保留证据约束下的 hedging、数字和来源。视觉样式参考 [vintage-latex](#) 与 [fiziko](#) 的科学文档设计机制，但未复制其模板或图形源码。

15 Source Ledger: 逐条溯源

A/B/C 仅表示来源类型与直接性。URL 可直接打开原始页面。So1 Coefficient Giving - About Us

Evidence B

official org page. Good Ventures is described as Coefficient's founding and most significant partnership; Coefficient serves as outsourced foundation staff.

<https://coefficientgiving.org/about-us/>

So2 Coefficient Giving - We're urgently scaling our work on AI and biosecurity

Evidence B

official org page. AI safety/security/field-building commitments: \$168m (2024), \$351m (2025), >\$1bn projected (2026); describes an outsized ecosystem role and institution-building.

<https://coefficientgiving.org/research/were-urgently-scaling-our-work-on-ai-and-biosecurity/>

So3 Coefficient Giving - Navigating Transformative AI

Evidence B

official strategy page. Strategy includes visibility/evaluations, safeguards, and talent/institutional capacity.

<https://coefficientgiving.org/funds/navigating-transformative-ai/>

So4 Good Ventures - Pattern Labs: Technological Risk Mitigation

Evidence A

official grant record. \$6.8m grant, awarded 02/2024, recommended by Open Philanthropy; Pattern Labs later became Irregular.

<https://www.goodventures.org/our-portfolio/grants/pattern-labs-technological-risk-mitigation/>

So5 Good Ventures - Grants database (historical)

Evidence A

official grant database. Includes Apollo Research startup funding \$1.53548m (2023), Redwood Research \$10.7m (2022) and \$5.3m (2023), among other records.

<https://www.goodventures.org/our-portfolio/grants-database/page/109/?limit=all>

So6 Good Ventures - Grants database (2024 page)

Evidence A

official grant database. Includes Apollo Research general support \$2.1787m (2024).

<https://www.goodventures.org/our-portfolio/grants-database/page/17/?limit=10000>

So7 ProPublica Nonprofit Explorer - Good Ventures Foundation

Evidence A

IRS-derived filing database. FY2025 assets about \$10.108bn and net assets about \$9.932bn; public summary does not establish the donated Anthropic stake is held by this legal entity.

<https://projects.propublica.org/nonprofits/organizations/461008520>

So8 Forbes - Inside a billionaire couple's plan to give away their fortune

Evidence C

secondary journalism. Reports Anthropic stake was moved into a nonprofit vehicle; does not identify Good Ventures Foundation as the holder.

<https://www.forbes.com/sites/phoebeliu/2025/11/07/cari-tuna-billionaire-open-philanthropy-facebook/>

So9 Forbes - Reranking the world's billionaires by wealth and altruism

Evidence C

secondary journalism. Reports donated early Anthropic investment and estimated stake below 0.8%; ownership destination remains unresolved publicly.

<https://www.forbes.com/sites/mattdurot/2026/04/20/reranking-the-worlds-billionaires-by-wealth-and-altruism/>

S10 MATS - Transparency

Evidence A

official transparency page. 2025 itemized cash donations used for donor-share and HHI calculations.

<https://www.matsprogram.org/transparency>

S11 MATS - Alumni: Marius Hobbhahn

Evidence B

official bio. MATS alumnus; later CEO/director of Apollo Research.

<https://www.matsprogram.org/alumni/marius-hobbhahn>

S12 MATS - Alumni: Joseph Bloom

Evidence B

official bio. MATS alumnus; later Head of Model Transparency at UK AISI.

<https://www.matsprogram.org/alumni/joseph-bloom>

S13 MATS - Alumni: Quentin Feuillede-Montixi

Evidence B

official bio. After MATS, worked as METR contractor on pre-release model evaluation.

<https://www.matsprogram.org/alumni/quentin-feuillede-montixi>

S14 MATS - Summer 2026 Program

Evidence B

official program page. Post-program pathways explicitly include frontier labs, evaluators, AI-safety funders, and new alignment organizations.

<https://www.matsprogram.org/program/summer-2026>

S15 Future of Life Institute - Finances

Evidence A

audited/official financial disclosure. 2025 total expenditures \$46.192m; grants \$26.724m; media and outreach \$5.795m.

<https://futureoflife.org/about-us/finances/>

S16 Future of Life Institute - Digital Media Accelerator

Evidence B

official program page. Supports creators raising awareness/understanding of AI safety/risk; states plan to spend up to \$100k/month across projects.

<https://futureoflife.org/project/digital-media-accelerator/>

S17 Future of Life Institute - Multistakeholder Engagement grants

Evidence B

official grant program. Up to \$5m program for public/stakeholder education, outreach, grassroots organization and movement-building.

<https://futureoflife.org/grant-program/multistakeholder-engagement-for-safe-and-prosperous-ai/>

S18 Future of Life Institute - Protect What's Human campaign

Evidence B

official press release. Campaign announced as potentially spending up to \$8m on AI-regulation advocacy.

<https://futureoflife.org/press-release/fli-launches-multimillion-dollar-nationwide-ai-regulation-campaign/>

S19 Tarbell Center - Ethics and Standards

Evidence B

official policy page. Says majority of 2025 funding originated with Coefficient; also states donors have no editorial or personnel approval rights.

<https://www.tarbellcenter.org/ethics-and-standards>

S20 Tarbell Center - Fellowship

Evidence B

official program page. Nine-month newsroom placement, AI journalism training, Bay Area summit, and stipends.

<https://www.tarbellcenter.org/fellowship>

S21 Tarbell Center - Kai Williams

Evidence B

official bio. Illustrative overlap: prior MATS AI-safety research experience and later Tarbell journalism fellowship.

<https://www.tarbellcenter.org/team/kai-williams>

S22 METR - Funding Update (Aug 2026)

Evidence B

official funding disclosure. Reports roughly \$71m in commitments over prior six months from multiple donors.

<https://metr.org/blog/2026-08-14-funding-update/>

S23 OpenAI - Third-party cyber evaluations involving OpenAI models

Evidence A

official incident disclosure. Irregular CTF environment was misconfigured to allow internet access; incident was not a sophisticated sandbox escape or zero-day.

<https://openai.com/index/third-party-cyber-evaluations-involving-openai-models/>

S24 Anthropic - Investigating incidents in cybersecurity evaluations

Evidence A

official incident disclosure. Large retrospective review; incidents involved third-party evaluation environment and unintended internet access.

<https://www.anthropic.com/news/investigating-incidents-cybersecurity-evals>

S25 Anthropic - Alignment assessment of cybersecurity incidents

Evidence A

official technical analysis. Four incidents shared one evaluation partner; prompts lacked explicit scope boundaries; includes scope-reminder, realism, and safe-exit interventions.

<https://www.anthropic.com/research/alignment-assessment-cybersecurity-incidents>

S26 Anthropic - Improving alignment and security efforts

Evidence A

official practices update. Recommends explicit in/out-of-scope targets, permitted actions, network boundaries, and real-time monitoring.

<https://www.anthropic.com/news/improving-alignment-security-efforts>

S27 Irregular - Addressing recent incidents

Evidence A

official postmortem. States disclosures arose from same underlying issue, incidents were rare, and better implementation of existing safeguards could prevent most incidents of this kind.

<https://www.irregular.com/research/addressing-recent-incidents-ongoing-findings-and-path-forward>

S28 OpenAI - Hugging Face incident and the road ahead

Evidence A

official incident disclosure. Separate incident involving containment circumvention; should not be conflated with Irregular misconfiguration events.

<https://openai.com/index/hugging-face-incident-and-the-road-ahead/>

S29 OpenAI - Foundations for trustworthy third-party evaluations

Evidence A

official methods page. Evaluation results depend on model, environment and setup; supports provenance-aware evaluation practice.

<https://openai.com/index/trustworthy-third-party-evaluations-foundations/>

S30 OpenAI - GPT-4o System Card

Evidence A

official system card. Uses METR and Apollo Research as third-party evaluators.

<https://openai.com/index/gpt-4o-system-card/>

S31 OpenAI - o1 System Card

Evidence A

official system card. Uses METR and Apollo Research among external evaluators.

<https://openai.com/index/openai-o1-system-card/>

S32 OpenAI - GPT-5.6 Deployment Safety

Evidence A

official deployment safety report. Uses METR, Apollo Research, UK AISI, and Irregular across frontier-risk evaluations.

<https://deploymentsafety.openai.com/gpt-5-6>

S33 OpenAI - GPT-6 Astra Deployment Safety

Evidence A

official deployment safety report. Uses Apollo Research, UK AISI, and Irregular; Irregular cyber work described as sandboxed without public internet.

<https://deploymentsafety.openai.com/gpt-6-astra>

S34 Apollo Research - About

Evidence B

official org page. Says it ran evaluations for all major labs in 2025 and increasingly engages in research-agenda/policy work.

<https://www.apolloresearch.ai/about>

S35 Apollo Research - Norms, conflicts of interest, security, science and communication

Evidence B

official policy page. COI rules include no outcome-contingent compensation and recusal for relevant financial interests.

<https://www.apolloresearch.ai/blog/our-norms-coi-security-science-communication>

S36 Apollo Research - Becoming a PBC

Evidence B

official org announcement. Describes growing commercial market for frontier AI safety and intent to build out/shape that market.

<https://www.apolloresearch.ai/blog/apollo-research-is-becoming-a-pbc>

S37 Reuters Institute - Digital News Report 2026 Executive Summary

Evidence A

research report. Across 48 markets, social media/video are major news gateways; useful context for exposure-level framing.

<https://reutersinstitute.politics.ox.ac.uk/digital-news-report/2026/dnr-executive-summary>

S38 Nature Human Behaviour - Sharing without clicking on news in social media

Evidence A

peer-reviewed research. Large Facebook-post study reports roughly three-quarters of shared news links were shared without clicking.

<https://www.nature.com/articles/s41562-024-02067-4>

S39 Ecker et al. - The effects of subtle misinformation in news headlines

Evidence A

peer-reviewed research. Misleading headlines affect memory, inferential reasoning and behavioral intentions even when full articles are read.

<https://pubmed.ncbi.nlm.nih.gov/25347407/>

S40 Los Angeles Times - Is there really a 10% chance AI could kill us all?

Evidence C

news article. Case study: strong catastrophic framing in headline/deck; later includes optimistic/physical-constraint and regulatory-incentive counterarguments; author disclosure notes Tarbell support.

<https://www.latimes.com/business/story/2026-09-11/is-there-really-10-chance-ai-could-kill-us-all>

S41 Washington Post - For years they warned AI could kill all humans. Now people are listening Evidence C

secondary journalism. Independent reporting on donor-funded AI-safety pipelines in universities, government and media; also includes skeptical voices.

<https://www.washingtonpost.com/technology/2026/09/10/years-they-warned-ai-could-kill-all-humans-now-people-are-listening/>

S42 Reuters - “Going rogue” draws critics amid widening AI hacks Evidence C

secondary journalism. Reports criticism that anthropomorphic “going rogue” framing can distort interpretation of engineering incidents.

<https://www.reuters.com/technology/artificial-intelligence/going-rogue-draws-critics-amid-widening-ai-hacks-2026-08-05/>

Do1 Design reference - fiziko Evidence -

design reference. MetaPost scientific illustration library; used only as visual inspiration, not as evidence or copied report code.

<https://github.com/jemmybutton/fiziko>

Do2 Design reference - vintage-latex Evidence -

design reference. Collection of vintage scientific-paper LaTeX mechanisms; used only as visual inspiration, not as evidence or copied report code.

<https://github.com/Foadsf/vintage-latex>
